DOE Genesis Mission

DE-FOA-0003612

A Lab PI's Guide to Writing a Winning Proposal

Navigating the Science–AI Integration

from a National Laboratory Perspective

Total Program Funding:	\$293.76 million
Phase I Awards:	\$500K–\$750K (9 months)
Phase II Awards:	\$1.5M–\$3.75M (3 years)
Phase I Deadline:	April 28, 2026, 11:59 PM ET
Project Narrative:	5 pages, 11pt, 1-inch margins

Analysis based on DE-FOA-0003612 issued March 17, 2026.
Always verify against the official NOFO and any amendments.

Contents

1 Quick Reference **2**

2 What the Review Panel Is Actually Looking For **2**
2.1 The Five Criteria — In Descending Order of Weight . . . 2
2.1.1 Criterion 1 (Highest Weight): Scientific and/or Technical Merit and Impact . 2
2.1.2 Criterion 2: Technical Approach, Methods, and Feasibility . . . 3
2.1.3 Criterion 3: Team, Resources, and Management . . . 3
2.1.4 Criterion 4: Commercialization Potential (Applied Technology Only) . . . 3
2.1.5 Criterion 5: Budget and Cost-Effectiveness . . . 3
2.2 "AI Advantage" Is Not a Buzzword — It Is the Decision Gate . . . 4
2.3 The 5-Page Narrative: Recommended Structure . . . 4
2.4 Program Policy Factors (Post-Review Selection) . . . 5
2.5 Genesis Mission Integration Requirements . . . 5
2.6 Portrait of a Top-Scoring Phase I Proposal . . . 6

3 The Science–AI Question: A Candid Assessment **7**
3.1 The Uncomfortable Truth: This Is Genuinely an AI Call . . . 7
3.2 The Equally Important Truth: The Science Cannot Be Thin . . . 7
3.3 What This Means for Your Relationship with Your Program Manager . . . 7
3.4 The Right Way to Integrate Science and AI . . . 8
3.4.1 Start from the Scientific Bottleneck . . . 8
3.4.2 Choose AI Methods That Match the Bottleneck . . . 9
3.4.3 Design the Experiment So AI Advantage Is Measurable . . . 9
3.4.4 Build the Team Around the Integration . . . 9
3.5 What Your Lab Brings That Others Cannot . . . 10
3.6 Lab-Specific Administrative Considerations . . . 10

4 Tactical Checklist for Lab PIs **12**
4.1 Before Writing (Weeks 1–2) . . . 12
4.2 Writing the Proposal (Weeks 2–4) . . . 12
4.3 Final Checks (Week Before Deadline) . . . 13

1 Quick Reference

Item	Detail
Total Funding	$293.76 million
Phase I Award	$500K–$750K for 9 months
Phase II Award	3–5× Phase I (i.e., $1.5M–$3.75M) for 3 years
Phase I Deadline	April 28, 2026, 11:59 PM ET
Phase II LOI	April 28, 2026, 5:00 PM ET
Project Narrative	5 pages — 11pt, 1-inch margins
Team Requirement	Partners from ≥2 of: (1) National Lab/User Facility, (2) Industry, (3) IHE/Nonprofit
Instrument	Other Transaction Agreement (OTA) — milestone-based
Cost Share	Not required for non-profits/academia; 20% for industry R&D, 50% for demonstration
Lab Authorization	Written authorization from cognizant Agreement Officer required (for non-SC labs: may need to be submitted with the application)

2 What the Review Panel Is Actually Looking For

This section unpacks the five review criteria in their NOFO-specified order of importance. If you internalize nothing else from this guide, internalize the ranking: scientific and technical merit is weighted highest, but **AI advantage is evaluated as part of that top criterion, not as a separate add-on.** The two are fused by design.

2.1 The Five Criteria — In Descending Order of Weight

2.1.1 Criterion 1 (Highest Weight): Scientific and/or Technical Merit and Impact

For **fundamental research**, reviewers are asked:

a) What is the scientific innovation of the proposed research?

b) How well does the proposed research achieve AI advantage?

c) What is the likelihood of achieving valuable results?

d) How might results impact the direction, progress, and thinking in relevant fields?

e) How original and innovative is the proposed scientific approach or concept?

f) How well does the work articulate its contribution to the scientific community or mission application?

Notice that (a) and (b) are the first two questions. A proposal with brilliant science but no credible AI advantage story will not score well on the single most important criterion. Conversely,

a proposal with sophisticated AI but thin science will fail on (a), (c), and (d). The reviewers are being told, explicitly, that **both** must be present and **both** must be strong.

For **applied technology development**, reviewers additionally evaluate TRL advancement (e.g., TRL 3 → TRL 5), KPIs for mission applications, and alignment with focus area objectives.

2.1.2 Criterion 2: Technical Approach, Methods, and Feasibility

- Is the methodology logically sound and well justified?
- Does the research employ innovative concepts or methods?
- Are experimental design, theoretical models, and/or computational simulations robust and detailed?
- Have potential challenges and risks been identified with reasonable mitigation strategies?
- Is there a clear, realistic timeline for key activities and milestones?

For a lab PI, this is where your existing infrastructure is a major asset. You can describe specific beamlines, test reactors, computing allocations, and experimental pipelines that make the 9-month timeline credible. University teams have to promise access; you can demonstrate it.

2.1.3 Criterion 3: Team, Resources, and Management

- How well qualified is the team?
- Does the team possess necessary expertise across all relevant disciplines?
- Is there clear, well-coordinated collaboration with balanced responsibilities, synergistic contributions, and clearly defined roles?
- Are facilities, equipment, and resources — **including access to data and AI models** — adequate?

Note the last bullet: **access to data and AI models** is called out explicitly. Reviewers are being told to check whether your team actually has the AI capability it claims. A named co-PI with a track record in scientific ML, or a partnership with an AI company that brings specific models, is far more credible than "we will hire a postdoc."

2.1.4 Criterion 4: Commercialization Potential (Applied Technology Only)

- Credible technology transfer/commercialization strategy
- Identified market opportunities and preliminary techno-economic analysis
- Industry adoption potential and barriers addressed

2.1.5 Criterion 5: Budget and Cost-Effectiveness

- Reasonable, well-justified budget
- Cost share meets requirements and reflects strong commitment
- Efficient use of resources and good value for federal investment; fostering collaboration with DOE/NNSA National Laboratories, academia, and industry

2.2 "AI Advantage" Is Not a Buzzword — It Is the Decision Gate

The term **"AI advantage"** is the central organizing concept of this entire NOFO. The NOFO defines Phase I success as:

> *Demonstrating a clear, tangible research workflow that incorporates AI with concrete evaluation of the potential for AI advantage.*

This is not asking you to *use* AI. It is asking you to *prove that AI helps*. That distinction matters enormously. The NOFO gives explicit examples of what constitutes demonstrated AI advantage:

1. **Increased predictive power or scientific insight** from appropriately-curated data
2. **More tightly coupling data and experiments** to validate hypotheses
3. **Building new models and analyzing their impact** on discovery speedup
4. **Identifying scaling metrics** that show how performance improves with more data or computing resources
5. **Improving and speeding up experimental workflows** (through automation or AI-informed parameters)
6. **Other proposed metrics** the team would like to be considered

> **Key Insight:** *Your Phase I go/no-go decision is fundamentally: did AI make the science measurably better? You need to define "better" quantitatively before you start, establish a baseline, and show the delta. Proposals that cannot articulate this before-and-after comparison in concrete terms will not survive review.*

2.3 The 5-Page Narrative: Recommended Structure

Five pages is brutally short. The NOFO recommends the following allocation:
Additionally, for **all** applications the narrative must:

- Clearly specify which aspects of realistic scientific or engineering workflows can be accelerated or improved by the proposed AI capabilities, and why AI is particularly promising compared to other potential solutions.
- Describe the **full pipeline**: data generation/curation → experiments/observation → AI model training/fine-tuning/inference → integration into a capability for science or engineering.
- If data is sensitive/proprietary, explain how data preparation and AI training will protect sensitive aspects while maximizing the value of the resulting AI capability.

> **Strategic Move:** As a lab PI, your "Data Sources & Models" half-page is potentially your strongest section. You have access to facility data, instrument archives, and simulation databases that university teams can only dream about. Make that access concrete: name the datasets, give approximate sizes, describe their provenance. Reviewers evaluating Criterion 3 will be checking whether your data story is real.

Section	Pages	What to Cover
Background/Introduction	∼1	Importance, relevance, which aspect of the chosen challenge you address. Position vs. other work. Novel aspects. Literature.
Project Objectives	∼0.5	Specific objectives. Alignment with focus topic. How they lead to AI advantage.
Proposed Research & Methods	∼1.5	Clear 9-month research plan. Activities, methods, technical detail. PI responsibility and budget per activity. Distinguish from existing funded work.
Milestones in Nine Months	∼1	Clearly defined and measurable milestones.
Data Sources & Models	∼0.5	Existing/new data sources. Your access to the data. AI models/frameworks to be used.
Decision Gate Metrics	∼0.5	Specific metrics for Phase I go/no-go. Scaling behavior preferred. Quantitative statements.

2.4 Program Policy Factors (Post-Review Selection)

After merit review, the Selection Official also weighs:

- Relevance to DOE priorities
- Appropriate balance across DOE programs
- Performance under current awards
- **Strengthening partnerships between DOE/NNSA labs, universities, non-profits, and industry**
- **Commitment to sharing results of research**

The partnership and sharing factors favor lab-led proposals that bring in external partners rather than keeping work in-house. A lab PI who recruits a strong university group and an industry partner — rather than staffing everything with lab FTEs — is signaling exactly what the Selection Official wants to see.

2.5 Genesis Mission Integration Requirements

Every team is expected to:

- Participate in coordination activities across the Genesis Mission
- Host data, models, and artifacts on the platform with metadata for discoverability and reuse
- Take advantage of the American Science Cloud (AmSC) and the Transformational AI Models Consortium (ModCon) when practical
- Build upon general-purpose capabilities produced by other Genesis Mission teams, creating domain-specific capabilities in a plug-and-play fashion

This is not boilerplate. The OTA instrument is milestone-based, and integration milestones will likely be part of your agreement. Proposals that describe concrete plans for contributing models, datasets, or workflows to the Genesis ecosystem — not just consuming resources from it — will be more competitive.

2.6 Portrait of a Top-Scoring Phase I Proposal

Based on the full NOFO analysis, a proposal that scores at the top will combine the following:

A scientific problem where AI is the rate-limiting breakthrough.
Not AI as a convenience or acceleration of something you could do anyway with more postdocs. The problem should be one where the dimensionality, nonlinearity, data volume, or speed requirement makes conventional approaches fundamentally inadequate. Reviewers can tell the difference between "AI would be nice" and "AI is the only realistic path."

A working prototype, not a research plan.
Phase I asks for a 9-month *demonstration*. The winning proposals will describe specific software, specific data, specific experiments, and a specific evaluation. "We will explore whether…" is the language of a proposal that will not be funded. "We will build X, train it on Y, validate against Z, and measure AI advantage as the ratio of…" is the language that wins.

Quantitative decision-gate metrics with baselines.
The NOFO explicitly asks for scaling behavior — performance that improves with more data or compute. This is a strong signal that reviewers will value neural scaling law-style analyses. But any quantitative comparison against a well-defined baseline will work: prediction accuracy vs. expert judgment, time-to-result vs. conventional workflow, or design-space coverage vs. grid search.

A data story that is already credible.
You either have the data now, or you have a concrete plan to generate it in months 1–3. "We will work with our partners to identify suitable datasets" is a proposal that has not done its homework. Name the datasets. Give sizes. Describe quality.

A team where every partner is essential.
The NOFO requires partners from ≥ 2 of the 3 categories and asks reviewers to evaluate "synergistic contributions" and "balanced responsibilities." A team where the university partner contributes a named AI researcher with a relevant publication record, and the industry partner provides deployment pathways or proprietary data, is far more credible than one where partners were recruited for eligibility compliance.

A visible path from Phase I to Phase II.
Phase I is $500–750K over 9 months, but reviewers are thinking about whether this approach scales to a 3-year, multi-million-dollar effort. Your Phase I proposal should make the Phase II vision legible even though you are not writing a Phase II proposal yet.

A concrete Genesis Mission integration plan.
How will your models, data, and workflows be contributed to the American Science Cloud? How will you engage with ModCon? How will other teams be able to build on your work? This is not an afterthought — it is part of the mission.

3 The Science–AI Question: A Candid Assessment

If you are a lab PI whose career is built on domain science — materials, biology, plasma physics, nuclear engineering, environmental systems — you may be hearing two conflicting messages. Your program manager says: "Focus on the science. That is what we fund." This NOFO says: "Demonstrate AI advantage." This section addresses the tension head-on, without sugarcoating either side.

3.1 The Uncomfortable Truth: This Is Genuinely an AI Call

Let us start with what this NOFO is *not*. It is not a traditional SC research call with an AI-friendly paragraph tucked into the supplementary information. It is not BER, BES, FES, or NE soliciting their usual portfolio and accepting AI-flavored proposals. The Genesis Mission is a $294 million, White House-priority initiative whose explicit purpose is to transform how DOE science is done through AI. The entire funding structure — OTA instrument, milestone-based payments, 9-month proof-of-concept phases, integration with a national AI platform — is designed around AI delivery, not traditional research timelines.

If you read the topic area descriptions, every single one follows the same template: **Challenge** (the science), **AI Solution** (how AI addresses it), **Justification** (DOE's unique position), **National Impact**. The AI Solution section is not decorative. It specifies the *kinds* of AI expected: agentic AI, digital twins, foundation models, physics-informed neural networks, surrogate models, autonomous experimental systems. These are not generic buzzwords — they are technical requirements.

> **Key Insight:** *A proposal that treats AI as a minor methodological detail — the equivalent of "we will also run some machine learning" — will not be competitive. The reviewers are being asked, as their second question under the highest-weighted criterion, "How well does the proposed research achieve AI advantage?" A vague or perfunctory answer to that question sinks the proposal regardless of how strong the science is.*

3.2 The Equally Important Truth: The Science Cannot Be Thin

That said, this is *also* not an ASCR call asking for AI methods research that happens to touch a domain. The review criteria put "scientific innovation" as question (a) under Criterion 1, ahead of "AI advantage" at question (b). The topic areas are defined by scientific challenges — genotype-to-phenotype prediction, critical mineral separation chemistry, fusion materials qualification, subsurface energy systems — not by AI techniques. Reviewers from BER, BES, FES, and NE will be evaluating whether the proposed science is important, novel, and well-conceived. A proposal that showcases a clever transformer architecture but addresses a trivial or already-solved scientific problem will fail on (a), (c), and (d) of Criterion 1.

The NOFO's structure reveals its actual logic: **the science defines the problem; the AI defines the approach; the proposal must be strong on both.** This is not "science with an AI garnish" and it is not "AI with a science veneer." It is a genuine integration requirement.

3.3 What This Means for Your Relationship with Your Program Manager

Your BER, BES, FES, or NE program manager is not wrong when they say "focus on the science." The science *must* be excellent. But they may be underestimating how seriously the NOFO takes the AI component. Here are specific ways to have that conversation:

Argument 1: The funding structure tells the story.

This is an OTA with milestone-based payments, not a standard financial assistance award. Milestones will include AI deliverables — trained models, validated workflows, scaling analyses, contributions to the American Science Cloud. If your milestones are purely scientific with AI as a side note, you may not pass the milestone reviews.

Argument 2: The review criteria are explicit.

Criterion 1 asks about scientific innovation *and* AI advantage as co-equal top-level questions. Criterion 3 specifically asks whether the team has adequate "access to data and AI models." Reviewers are being instructed to evaluate AI capability as a first-order concern, not a nice-to-have.

Argument 3: ASCR is co-funding across all topic areas.

The NOFO states that ASCR intends to identify co-funding opportunities across every partner office's focus areas. This means the AI component is not just tolerated — it is independently funded. Your program manager's science dollars and ASCR's AI dollars are designed to work together.

Argument 4: Phase I is a proof-of-concept, and the proof is about AI.

The decision gate metric for Phase I is whether AI advantage has been demonstrated. A team that produces excellent science but cannot show that AI made a measurable difference will not advance to Phase II. The 9-month timeline is calibrated to this demonstration, not to a full scientific investigation.

Argument 5: This is where the money is going.

$294 million is substantial. If your program's science priorities align with a Genesis topic area, this is a major funding opportunity. But it requires genuine engagement with AI, not a pro-forma mention. The question is whether your group is willing to invest the effort to build real AI capability, not whether you can write the word "AI" convincingly.

3.4 The Right Way to Integrate Science and AI

The proposals that will score highest are those where the science and AI are **genuinely inseparable** — where you cannot do the science without the AI, and the AI has no purpose without the science. Here is how to achieve that integration in practice:

3.4.1 Start from the Scientific Bottleneck

Identify the specific step in your research workflow where progress is currently blocked or unacceptably slow. Maybe it is searching a combinatorial materials space. Maybe it is predicting emergent properties from genomic data. Maybe it is optimizing plasma control in real time. The bottleneck must be real and documented — not hypothetical.

Then ask: *why* is this bottleneck hard? If the answer involves high dimensionality, nonlinear interactions, sparse data in a large space, the need for real-time inference, or integration across heterogeneous data types, you have a natural case for AI. If the answer is "we need more beam time" or "we need better reagents," AI is probably not the right tool and this is not the right call.

3.4.2 Choose AI Methods That Match the Bottleneck

The NOFO names specific AI approaches across different topic areas: physics-informed neural networks, digital twins, agentic AI for autonomous experiments, foundation models for scientific domains, surrogate models for expensive simulations. These are not interchangeable. A proposal that matches the right AI technique to the right scientific bottleneck — and explains *why* that match works — demonstrates technical depth that reviewers will reward.

> **Caution:** Do not propose to "apply large language models" to a problem that is fundamentally about physical simulation, or to "build a digital twin" of a system where the physics is already well-modeled and the real issue is data scarcity. Mismatched AI methods signal that the team does not understand the AI landscape well enough to execute.

3.4.3 Design the Experiment So AI Advantage Is Measurable

This is where many proposals will fail. You need:

1. **A baseline.** What does the current workflow produce without AI? How fast? How accurate? At what cost?
2. **An AI-enhanced workflow.** Same scientific objective, but with AI integrated at the bottleneck step.
3. **A fair comparison.** Same data, same evaluation criteria, same time horizon. The AI advantage must be attributable to the AI, not to other improvements you made simultaneously.
4. **A scaling analysis.** The NOFO explicitly values metrics showing that performance improves with more data or compute. This is not accidental — it signals that DOE is thinking about foundation models and wants to see evidence that your approach will get better as Genesis Mission resources grow.

3.4.4 Build the Team Around the Integration

The worst team structure for this call is a science group and an AI group working in parallel with a handoff in month 6. The best structure is one where AI researchers are embedded in the scientific workflow from day one — attending experimental planning meetings, helping design data collection protocols, iterating on model architecture based on scientific feedback.

Concretely, for a lab-led proposal, this might mean:

- A lab PI who owns the scientific problem and the experimental facilities
- A co-PI (at the lab, a university, or an AI company) who owns the AI methodology and has a publication record in scientific ML — not just in computer science venues, but in domain journals where ML was applied to real scientific problems
- A university partner who contributes students or postdocs doing the hands-on integration work, bridging the science and AI groups
- An industry partner who provides either (a) proprietary data that enriches training, (b) a deployment pathway for the resulting AI capability, or (c) AI infrastructure and engineering support

3.5 What Your Lab Brings That Others Cannot

As a lab PI, you have structural advantages in this competition that are worth being explicit about:

Data at scale.
You have decades of experimental data from user facilities, reactor operations, materials characterization campaigns, environmental monitoring networks, and large-scale simulations. Much of this data is not publicly available but could be curated into training sets for scientific AI. This is an enormous competitive advantage over university teams that would need to generate data from scratch.

Computational resources.
The NOFO notes that computing resources may be available from DOE/NNSA National Laboratories. You already have INCITE, ALCC, or Director's Discretionary allocations, or can request them. You know how to run at scale on Leadership Computing Facilities. University teams typically do not.

Integrated experimental–computational pipelines.
Many labs already operate autonomous or semi-autonomous experimental facilities. If you have an existing self-driving lab, a high-throughput screening capability, or a digital twin of a facility, you are starting from a much higher baseline than a team proposing to build these from scratch.

Institutional continuity.
A 9-month Phase I followed by a 3-year Phase II requires a stable team. Lab staff scientists provide continuity that a university team relying on graduating students cannot match.

Cross-disciplinary density.
A national laboratory has materials scientists, computational physicists, computer scientists, and AI researchers in the same institution, often in the same building. The NOFO values integrated teams — you can build one without the overhead of a multi-institutional collaboration agreement, and add external partners for the capabilities you lack.

> **Caution:** These advantages only matter if you use them. A lab proposal that reads like a university proposal — describing planned experiments rather than leveraging existing infrastructure — wastes its strongest card. Name the facility. Name the dataset. Name the computing allocation. Specificity is credibility.

3.6 Lab-Specific Administrative Considerations

- **Agreement Officer Authorization.** SC labs need written authorization from their cognizant Agreement Officer. For non-SC/NNSA labs (INL, LLNL, LANL, NLR, SNL, SRNL), this authorization **may need to be submitted with the application**, and failure to provide it may result in rejection before merit review. Start this process immediately.
- **Lab overhead rates** can make $500–750K feel small. Budget carefully. Consider whether some lab contributions can be structured as in-kind rather than funded effort.
- **One lead application per focus area per institution.** Coordinate internally. If another PI at your lab is targeting the same focus area, you need to resolve this before submission, not after.

- **OTA instrument.** This is not a standard cooperative agreement or grant. The OTA is milestone-based and uses a non-negotiable agreement template designed for the Genesis Mission. Understand what this means for IP, reporting, and payment timing before you commit.
- **Field Work Proposals (FWPs).** If your lab is selected, the work will be authorized through the DOE field work authorization system and performed under your M&O contract. No separate financial assistance award will be made. Ensure your lab's internal processes can accommodate the Genesis Mission's likely fast-start timeline.

4 Tactical Checklist for Lab PIs

4.1 Before Writing (Weeks 1–2)

- ☐ Identify your focus area and confirm no other PI at your lab is targeting the same one as lead
- ☐ Obtain (or initiate) written authorization from your cognizant Agreement Officer
- ☐ Confirm your PI is in a permanent or indefinitely extensible position
- ☐ Recruit partners from ≥2 of the 3 required categories (if lab-led, you need at least one of: university/nonprofit or industry)
- ☐ Identify your AI co-PI by name — someone with a publication record in scientific ML relevant to your domain
- ☐ Inventory your existing data assets: what datasets do you have, how large are they, what format, what quality?
- ☐ Define your scientific bottleneck and your baseline metric
- ☐ Contact the Genesis Mission Consortium partnership service if you need help finding partners
- ☐ Register for the March 26, 2026 informational webinar
- ☐ Confirm SAM.gov and Grants.gov registrations are current for your lab

4.2 Writing the Proposal (Weeks 2–4)

- ☐ **Abstract:** 1 page. Must include vision for science/energy initiatives AND description of how/why AI enhances workflows.
- ☐ **Title Page:** Use the exact format specified (lead institution, all partners, budget table, focus area). Use the Excel template from Grants.gov.
- ☐ **Project Narrative:** Exactly 5 pages. 11pt font. 1-inch margins. No URLs for supplementary info.
- ☐ In the narrative, lead with the science, integrate the AI throughout, and ensure every AI element connects to a scientific purpose.
- ☐ Include Buy America statement in first 2 pages if applicable (or state it does not apply).
- ☐ Define quantitative decision-gate metrics with baselines and AI-advantage targets.
- ☐ **Appendix 1:** Bibliography (no page limit, does not count against 5 pages).
- ☐ **Appendix 2:** Facilities & Resources (≤0.5 page). *This is where your lab advantage is most visible.* Name specific facilities, instruments, computing allocations, and data archives.
- ☐ **Appendix 3:** Equipment (≤0.5 page).
- ☐ **Appendix 4:** DMSP — NOT required for submission, but will be required at award negotiation. Address FAIR principles, open-source licensing, scientific software best practices.
- ☐ **Appendix 5:** Synergistic Activities (optional, 1 page per senior/key person).

- ☐ **Appendix 6:** Letters of Collaboration or Access (if needed — facility access confirmation letters are OK on Federal letterhead).
- ☐ **Appendix 7:** Other Attachments — institutional commitment letters, cost share documentation.
- ☐ **Biographical Sketches** via SciENcv (attached to Senior/Key Person Profile form, NOT in the narrative).
- ☐ **Current and Pending Support** via SciENcv. Ensure complete listing of all activities regardless of funding source.
- ☐ **Collaborator Template** (Excel) attached to Field 12.
- ☐ **AI disclosure:** Disclose use of AI tools in preparing the application (unless used only for editing).

4.3 Final Checks (Week Before Deadline)

- ☐ All collaborative applications have identical Project Narratives including title
- ☐ No sensitive PII in any document (no SSNs, birthdates, citizenship, marital status, home addresses)
- ☐ No letters of support from Federal officials or on Federal letterhead (access/facility confirmation letters are OK)
- ☐ Budget uses current negotiated indirect cost and fringe benefit rates
- ☐ For-profit partners include required cost share (20% R&D / 50% demonstration)
- ☐ Lab authorization letter obtained and filed (for non-SC labs: included with application)
- ☐ Project Narrative is a single machine-readable PDF, not scanned, no editable fields or signatures
- ☐ All file names $\leq$50 characters, all unique, using only permitted characters
- ☐ Submit well before April 28, 2026, 11:59 PM ET — Grants.gov recommends 24–48 hours early
- ☐ Verify receipt: you should receive a Grants.gov tracking number (GRANTXXXXXXXX) by email

---

Analysis based on DE-FOA-0003612 issued March 17, 2026.
Always verify against the official NOFO and any amendments.